

Exercise 7 Phil Szalay, Florian Schneider

January 30, 2020

```
In [70]: import numpy as np
import random
import matplotlib.pyplot as plt
from statistics import mean

In [4]: number_of_possible_actions = 10

In [168]: def run_experiment(epsilon):
    # q: array of true rewards for action a
    q = np.random.normal(0, 1, number_of_possible_actions)
    optimal_action = np.argmax(q)

    number_of_optimal_actions = np.zeros(1000)

    # rewards: array of arrays of rewards for action a for every iteration
    rewards = np.zeros((2000, 1000))

    for j in range(2000):
        # Q: array of estimated reward for action a
        # n: dictionary of number of action a was choosen
        Q = np.zeros(number_of_possible_actions)
        n = np.zeros(number_of_possible_actions)

        for i in range(1000):
            random_number = random.uniform(0, 1)

            # calc next action a
            if (random_number > epsilon):
                action = get_max_index(Q)
            else:
                random_index = random.randint(0, 9)
                action = random_index

            if (action == optimal_action):
                number_of_optimal_actions[i] += 1

        reward = bandit(action, q)
```

```

        rewards[j, i] = reward

        n[action] += 1
        Q[action] = Q[action] + 1 / n[action] * (reward - Q[action])

    # calc average Q for all actions
    reward_averages = np.mean(rewards, axis=0)
    number_of_optimal_actions_averages = number_of_optimal_actions / 2000

    return (reward_averages, number_of_optimal_actions_averages)

# returns draw from normal distribution with mean of Q[action]
def bandit(action, q):
    return np.random.normal(q[action], 1)

# returns index of max element, breaking ties randomly
def get_max_index(Q):
    return np.random.choice(np.flatnonzero(Q == Q.max()))

In [169]: # epsilon = 0.1
a = run_experiment(0.1)

# epsilon 0.01
b = run_experiment(0.01)

# epsilon = 0
c = run_experiment(0)

In [170]: # plot all three curves
plt.plot(np.arange(1, 1001), a[0], label='epsilon = 0.1')
plt.plot(np.arange(1, 1001), b[0], label='epsilon = 0.01')
plt.plot(np.arange(1, 1001), c[0], label='epsilon = 0.0')
plt.legend()

plt.xlabel('Steps')
plt.ylabel('Average reward')
plt.title('Average reward epsilon-Greedy for different epsilons')
plt.show()

# plot optimal choices
plt.plot(np.arange(1, 1001), a[1], label='epsilon = 0.1')
plt.plot(np.arange(1, 1001), b[1], label='epsilon = 0.01')
plt.plot(np.arange(1, 1001), c[1], label='epsilon = 0.0')
plt.legend()

plt.xlabel('Steps')
plt.ylabel('% Optimal action')
plt.title('Optimal action epsilon-Greedy for different epsilons')
plt.show()

```

